

UNIX basics and introduction to HPC

Larisa M. Soto

Computational and Statistical Genomics Laboratory

Department of human genetics

20-Septemer-2021

Mission : aims to deliver inter-disciplinary research programs and empower the use of data in health research and health care delivery

McGill.CA / MCGILL INITIATIVE IN COMPUTATIONAL MEDICINE

Contact



MiCoM McGill initiative in
Computational Medicine

McGill initiative in Computational Medicine
740, Dr. Penfield Avenue, Montreal, Quebec,
Canada, H3A 0G1
email: info-micm@mcgill.ca

[Signup](#) to our newsletter to receive the latest news

<https://www.mcgill.ca/micm>

Workshop outline

1

UNIX operating system
File system structure
File management
Built-in commands

2

Variables
Pipes, loops and conditional statements
Shell scripting
Text processing and pattern matching

3

What is HPC?
Compute Canada
Servers, clusters, and supercomputers
Submitting jobs to a cluster



Dennis Ritchie (standing) and Ken Thompson (seating) at Bell Labs, 1960's

UNIX Operating system

- General purpose
- Multiuser
- Interactive
- Resource-efficient
- Hierarchical file system
- Compatibility (files, devices, processes)
- Time sharing

The Single UNIX[®] Specification History & Timeline

| | | |
|------|---------------------------------|---|
| 1969 | The Beginning | The history of UNIX starts back in 1969, when Ken Thompson, Dennis Ritchie and others started working on the "little-used PDP-7 in a corner" at Bell Labs and what was to become UNIX. |
| 1971 | First Edition | It had an assembler for a PDP-11/20, file system, fork(), roff and ed. It was used for text processing of patent documents. |
| 1972 | First UNIX Installations | The first installations had 3 users, no memory protection, and a 500 KB disk. |
| 1973 | Fourth Edition | It was rewritten in C. This made it portable and changed the history of OS's. |
| 1975 | Sixth Edition | UNIX leaves home. Also widely known as Version 6, this is the first to be widely available outside of Bell Labs. The first BSD version (1.x) was derived from V6. |
| 1979 | Seventh Edition | It was an "improvement over all preceding and following Unices" [Bourne]. It had C, UUCP and the Bourne shell. It was ported to the VAX and the kernel was more than 40 Kilobytes (K). |
| 1980 | Xenix | Microsoft introduces Xenix. 32V and 4BSD introduced. |
| 1982 | System III | AT&T's UNIX System Group (USG) release System III, the first public release outside Bell Laboratories. SunOS 1.0 ships. HP-UX introduced. Ultrix-11 introduced. |
| 1983 | System V | Computer Research Group, UNIX System Group (USG) and a third group merge to become UNIX System Development Lab. AT&T announces UNIX System V, the first supported release. Installed base 45,000. |
| 1984 | 4.2BSD | University of California at Berkeley releases 4.2BSD, includes TCP/IP, new signals and much more. X/Open formed. |
| 1984 | SVR2 | System V Release 2 introduced. At this time there are 100,000 UNIX installations around the world. |
| 1986 | 4.3BSD | 4.3BSD released, including internet name server. SVID introduced. NFS shipped. AIX announced. Installed base 250,000. |
| 1987 | SVR3 | System V Release 3 including STREAMS, TLI, RFS. At this time there are 750,000 UNIX installations around the world. IRIX introduced. |
| 1988 | | POSIX.1 published. Open Software Foundation (OSF) and UNIX International (UI) formed. Ultrix 4.2 ships. |
| 1989 | | AT&T UNIX Software Operation formed in preparation for spinoff of UNIX development group. Motif 1.0 ships. |
| 1989 | SVR4 | UNIX System V Release 4 ships, unifying System V, BSD and Xenix. Installed base 1.2 million. |
| 1990 | XPG3 | X/Open launches XPG3 Brand. OSF/1 debuts. Plan 9 from Bell Labs ships. |

"The Single UNIX Specification brings all the benefits of a single standard operating system, namely application and information portability, scalability, flexibility and freedom of choice for customers"

Allen Brown, President and CEO, The Open Group

The Story of the License Plate...

In 1983 Digital Equipment Corporation (DEC) was ramping up their engineering group to create and ship their first UNIX system product. One of the stalwarts of the group was Armando P. Stettner. Armando was a UNIX system devotee. He lived and he breathed the UNIX system. When he got his new car, it was natural that he got vanity license plates that said "UNIX" on them. And it only made it better that the state motto of New Hampshire was "Live Free or Die".

Armando often got requests from people along the lines of "When will we be able to get our UNIX system license directly from DEC?" And Armando kept saying "Real Soon Now" (RSN). Armando was going to a conference and he was dreading having to say "RSN" many more times, so he had a bright idea. Armando went prepared to give out "UNIX licenses". On stage, when the question came up, "When will we be able to get our UNIX system license directly from DEC?", Armando yelled "Right Now!" and produced facsimiles of his license plate, holding it up for all to see. It was an almost perfect likeness of his license plate, with the trademark "UNIX" in



the middle of it, but instead of having "Live Free or Die" across the bottom of the plate (as in the real case), it had it across the top. Across the bottom was the trademark acknowledgement.

DEC made UNIX license plates up in small numbers and handed them out at events. They usually ran out. The demand for the license plates never did abate. People saw them on an office wall, or heard about them somewhere, and wanted one of their own.

Armando left the state for the sunny climes of California, and had taken his car and license plate with him. Or so many people thought.

In 1989 Jon 'maddog' Hall was purchasing a new car, a Jeep Wrangler. And of course the license plate had to be relevant. So Jon, a long time DEC employee and UNIX system guru, submitted his application with many variations and the clerk said "I think we can give you your first choice..." and gave him the temporary paper plates (to be used on the car until the metal plates were manufactured) with "UNIX" on them. And so it has been ever since. Jon's Jeep has been the holder of the UNIX license plate.

The Open Group thanks Jon "maddog" Hall for sharing the story of the UNIX license plate.

The UNIX Brand

The UNIX Brand is used to identify products that have been certified as conforming to the Single UNIX Specification, initially UNIX 93, followed subsequently by UNIX 95, UNIX 98 and now UNIX 03.

| | | |
|-----------|---|--|
| 1991 | | UNIX System Laboratories (USL) becomes a company - majority-owned by AT&T. Linus Torvalds commences Linux development. Solaris 1.0 debuts. |
| 1992 | SVR4.2 | USL releases UNIX System V Release 4.2 (Destiny). October - XPG4 Brand launched by X/Open. December 22nd - Novell announces intent to acquire USL. Solaris 2.0 and HP-UX 9.0 ship. |
| 1993 | 4.4BSD | 4.4BSD the final release from Berkeley. June 16 - Novell acquires USL. |
| Late 1993 | SVR4.2MP | Novell decides to get out of the UNIX business. Rather than sell the business as a single entity, Novell transfers the rights to the UNIX trademark and the specification to X/Open Company. COSE Initiative delivers "Spec 1170" to X/Open for fasttrack. In December Novell ships SVR4.2MP, the final USL OEM release of System V. |
| 1994 | Single UNIX Specification | BSD 4.4-Lite eliminated all code claimed to infringe on USL/Novell. As the owner of the UNIX trademark, X/Open introduces the Single UNIX Specification (formerly Spec 1170) which separates the UNIX trademark from any actual code stream itself, thus allowing multiple implementations. |
| 1995 | UNIX 95 | X/Open introduces the UNIX 95 branding program for implementations of the Single UNIX Specification. Novell sells UnixWare business to SCO. Digital UNIX introduced. UnixWare 2.0 ships. OpenServer 5.0 debuts. |
| 1996 | | The Open Group forms as a merger of the Open Software Foundation (OSF) and X/Open. UnixWare 2.1, HP-UX 10.20 and IRIX 6.2 ship. |
| 1997 | Single UNIX Specification, Version 2 | The Open Group introduces Version 2 of the Single UNIX Specification, including support for realtime, threads and 64-bit and larger processors. The specification is made freely available on the web. IRIX 6.4, AIX 4.3 and HP-UX 11 ship. |
| 1998 | UNIX 98 | The Open Group introduces the UNIX 98 family of brands, including Base, Workstation and Server. First UNIX 98 registered products shipped by Sun, IBM and NCR. The Open Source movement starts to take off with announcements from Netscape and IBM. UnixWare 7 and IRIX 6.5 ship. |
| 1999 | UNIX at 30 | The UNIX system reaches thirty. Solaris 7 ships. Linux 2.2 kernel released. The Open Group and the IEEE commence joint development of a revision to POSIX and the Single UNIX Specification. First LinuxWorld conferences. Dot com fever on the stock markets. Tru64 UNIX ships. |
| 2001 | Single UNIX Specification, Version 3 | Version 3 of the Single UNIX Specification unites IEEE POSIX, The Open Group and the industry efforts. Linux 2.4 kernel released. The value of procurements of open systems referencing the UNIX brand exceeds \$55 billion. AIX 5L ships. |
| 2003 | ISO/IEC 9945 | The core volumes of Version 3 of the Single UNIX Specification are approved as an international standard. "Westwood" test suites shipped for UNIX 03 brand. Solaris 9.0 E ships. Linux 2.6 kernel released. |

Trademark Information: Plan 9 is a trademark and OSF/1 and The Open Group are registered trademarks of The Open Group in the United States and other countries. Linux is a registered trademark of Linus Torvalds. All other trademarks are the property of their respective owners. Copyright © 2004 The Open Group. All Rights Reserved.

UNIX is a registered trademark of The Open Group in the United States and other countries.

THE Open GROUP
Making standards work™

How it started



How it's going



Parts

The kernel

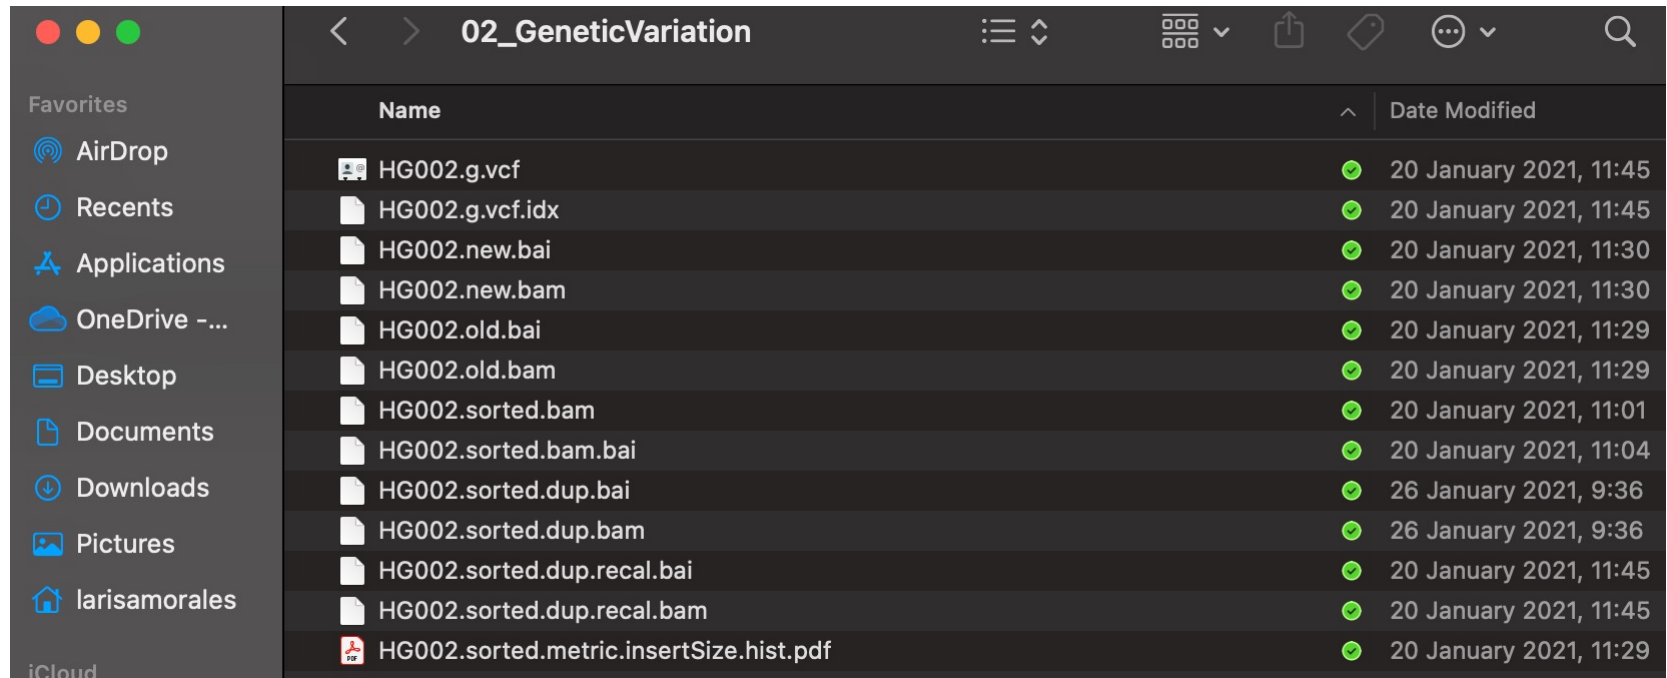
- The core of the operating system
- Manages processes' time, memory and resources
- Handles file storage
- Responds to system calls

The shell

- Let's the user communicate with the kernel
- Command line interpreter (CLI)
- Executes the instructions requested by the user (commands)

In UNIX
everything is a file or a
process!

Files



Files

```
total 132464
-rw-r--r--@ 1 larisamoraless  staff      75629 Jan 20  2021 HG002.g.vcf
-rw-r--r--@ 1 larisamoraless  staff     44048 Jan 20  2021 HG002.g.vcf.idx
-rw-r--r--@ 1 larisamoraless  staff     21584 Jan 20  2021 HG002.new.bai
-rw-r--r--@ 1 larisamoraless  staff     78169 Jan 20  2021 HG002.new.bam
-rw-r--r--@ 1 larisamoraless  staff     21584 Jan 20  2021 HG002.old.bai
-rw-r--r--@ 1 larisamoraless  staff     76424 Jan 20  2021 HG002.old.bam
-rw-r-----@ 1 larisamoraless  staff  17803945 Jan 20  2021 HG002.sorted.bam
-rw-r----- 1 larisamoraless  staff   1350440 Jan 20  2021 HG002.sorted.bam.bai
-rw-r--r--@ 1 larisamoraless  staff   1350440 Jan 26  2021 HG002.sorted.dup.bai
-rw-r--r--@ 1 larisamoraless  staff  17956229 Jan 26  2021 HG002.sorted.dup.bam
-rw-r--r--@ 1 larisamoraless  staff   1350440 Jan 20  2021 HG002.sorted.dup.recal.bai
-rw-r--r--@ 1 larisamoraless  staff  24793574 Jan 20  2021 HG002.sorted.dup.recal.bam
-rw-r--r--@ 1 larisamoraless  staff     32357 Jan 20  2021 HG002.sorted.metric.insertSize.hist.pdf
```

Processes

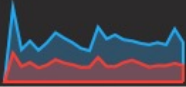
| Process Name | % CPU | CPU Time | Threads | Idle Wake Ups | % GPU | GPU Time | PID | User |
|-----------------------|-------|-------------|---------|---------------|-------|------------|-------|---------------|
| zoom.us | 1.7 | 52:25.01 | 26 | 262 | 0.0 | 0.00 | 47243 | larisamorales |
| XprotectService | 0.0 | 0.08 | 2 | 0 | 0.0 | 0.00 | 50628 | root |
| XprotectService | 0.0 | 0.06 | 2 | 0 | 0.0 | 0.00 | 50596 | larisamorales |
| xartstorageremoted | 0.0 | 2.62 | 2 | 0 | 0.0 | 0.00 | 235 | root |
| WirelessRadioManagerd | 0.0 | 0.13 | 2 | 1 | 0.0 | 0.00 | 50460 | root |
| WindowServer | 24.6 | 12:53:27.49 | 14 | 46 | 2.1 | 3:43:30.81 | 127 | _windowserver |
| wifianalyticsd | 0.0 | 2.60 | 2 | 0 | 0.0 | 0.00 | 293 | root |
| Wi-Fi | 0.0 | 5:28.75 | 4 | 0 | 0.0 | 0.00 | 369 | larisamorales |

System: 11.10%

User: 14.32%

Idle: 74.58%

CPU LOAD



Threads: 1,882

Processes: 384

Processes

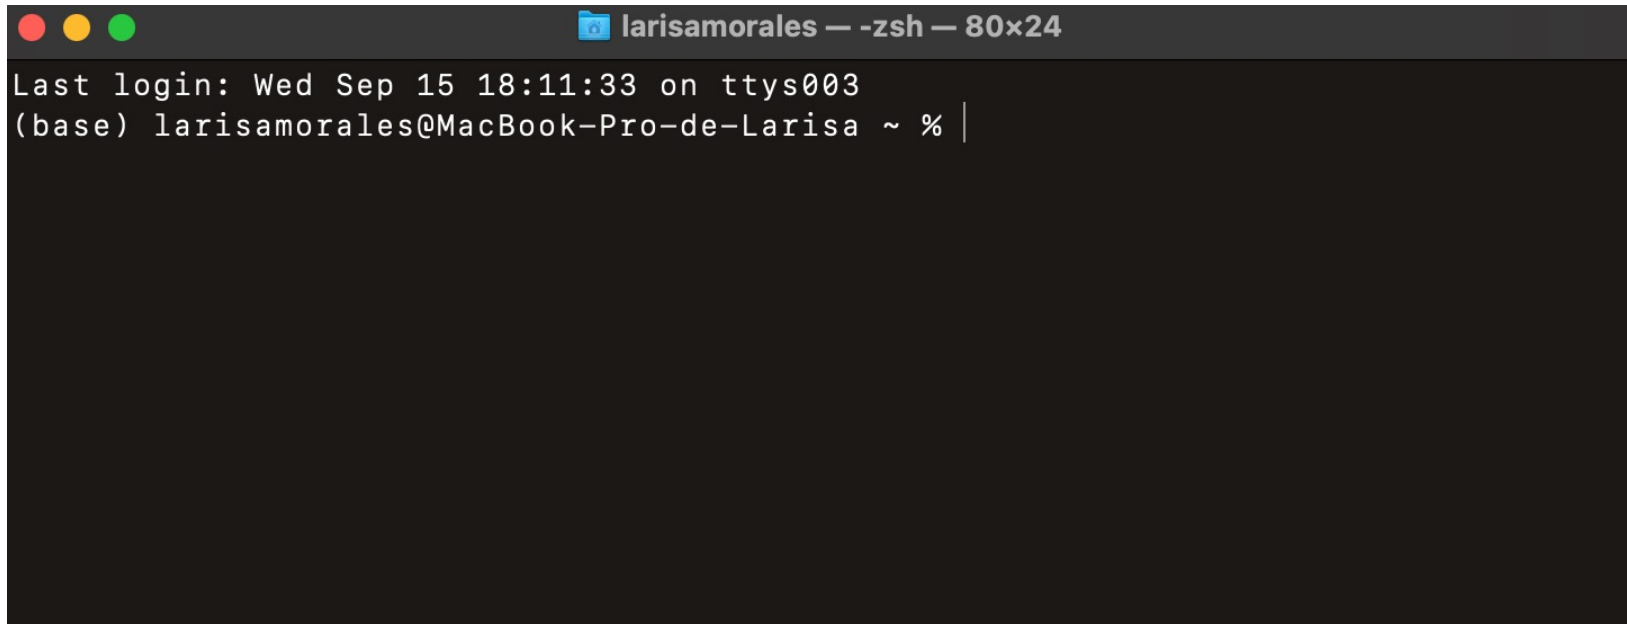
```

1  [|||||||||||||||||||||||||||||] 72.0%] Tasks: 384, 959 thr; 1 running
2  [|||||] 15.3%] Load average: 7.22 4.38 3.57
3  [|||||||||||||||||] 44.7%] Uptime: 25 days, 20:36:19
4  [|||||] 11.3%]
Mem[|||||||||||||||||||||||||] 4.26G/8.00G]
Swp[|||||||||||||||||] 3.03G/4.00G]

```

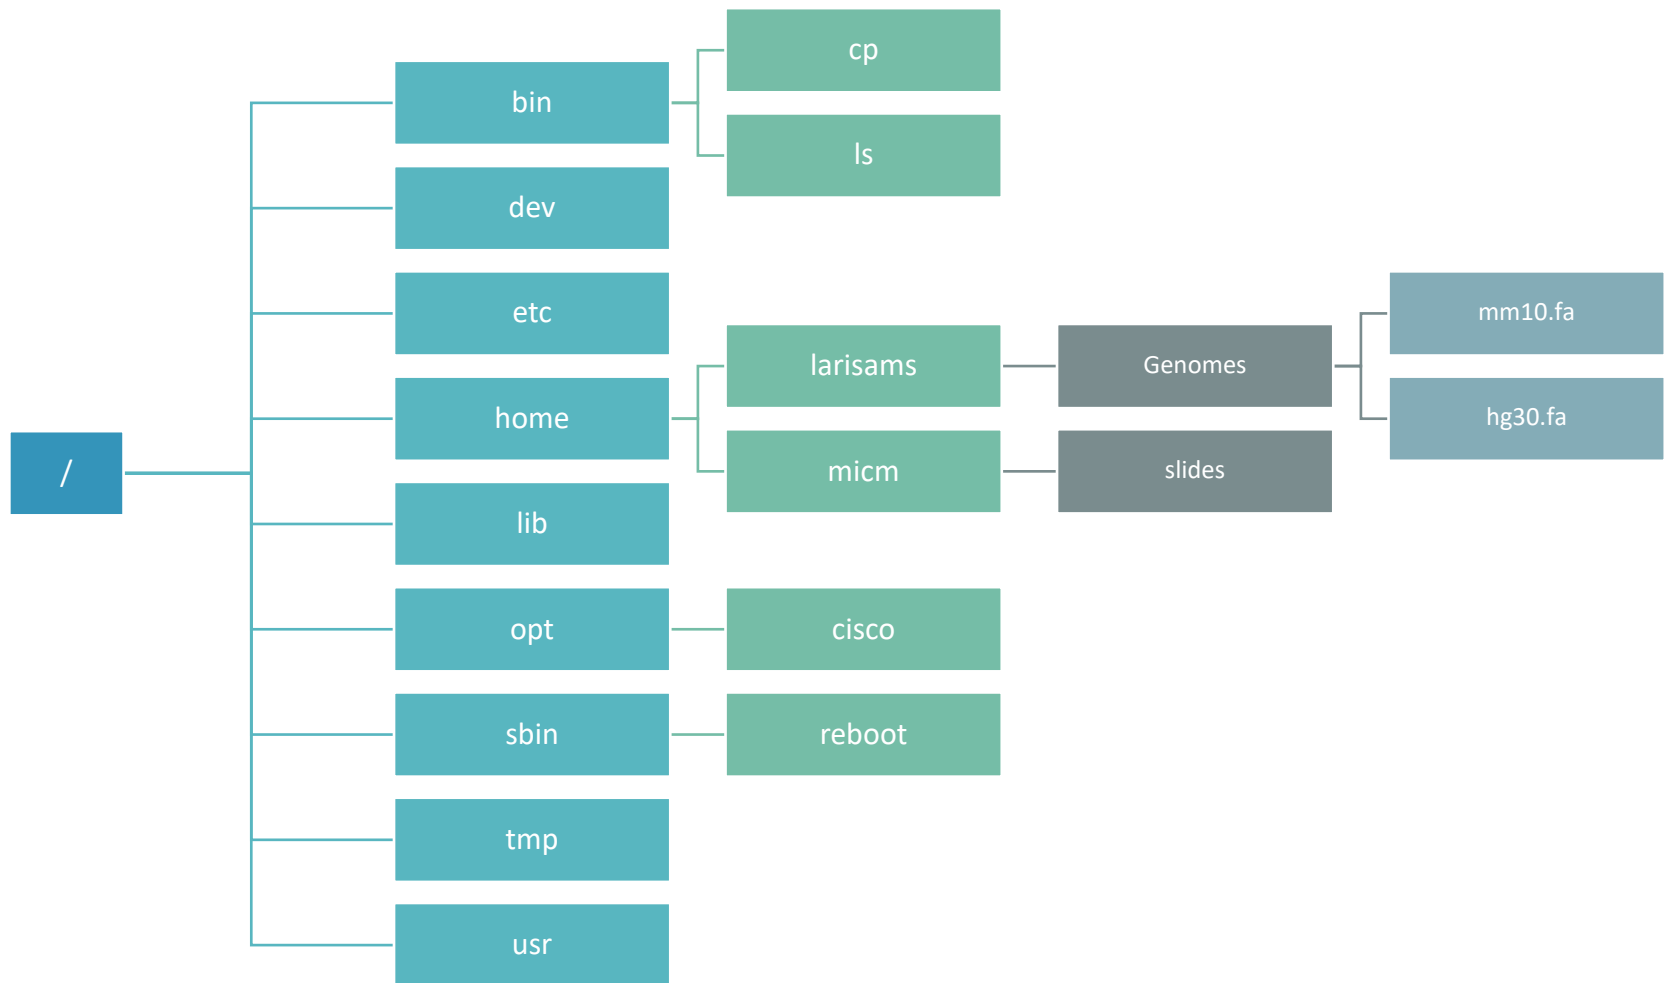
| PID | USER | PRI | NI | VIRT | RES | S | CPU% | MEM% | TIME+ | Command |
|-------|-----------|-----|----|-------|-------|---|------|------|----------|---|
| 45042 | larisamor | 17 | 0 | 8903M | 290M | ? | 4.2 | 3.5 | 13:30.28 | /Applications/Microsoft PowerPoint.app/Con |
| 51041 | larisamor | 17 | 0 | 4267M | 9768 | ? | 3.2 | 0.1 | 0:00.31 | /usr/sbin/screencapture -pdi -z cmd-shift- |
| 29417 | larisamor | 24 | 0 | 5682M | 161M | ? | 2.9 | 2.0 | 31:04.51 | /System/Applications/Utilities/Terminal.ap |
| 50407 | larisamor | 17 | 0 | 4301M | 8344 | ? | 1.6 | 0.1 | 1:17.94 | /System/Library/Frameworks/VideoToolbox.fr |
| 48033 | larisamor | 17 | 0 | 5485M | 53308 | ? | 1.3 | 0.6 | 5:22.14 | /Applications/Spotify.app/Contents/MacOS/S |
| 346 | larisamor | 17 | 0 | 5008M | 17224 | ? | 0.8 | 0.2 | 43:14.67 | /System/Library/CoreServices/ControlCenter |
| 384 | larisamor | 17 | 0 | 5201M | 14444 | ? | 0.7 | 0.2 | 1h04:46 | /System/Library/CoreServices/NotificationC |
| 51042 | larisamor | 17 | 0 | 4807M | 22880 | ? | 0.4 | 0.3 | 0:00.25 | /System/Library/CoreServices/screencapture |
| 46108 | larisamor | 17 | 0 | 6160M | 5604 | ? | 0.2 | 0.1 | 16:17.32 | /Applications/RStudio.app/Contents/MacOS/r |
| 415 | larisamor | 24 | 0 | 4325M | 11580 | ? | 0.2 | 0.1 | 10:47.17 | /usr/libexec/sharingd |
| 536 | larisamor | 17 | 0 | 4308M | 5124 | ? | 0.2 | 0.1 | 0:59.37 | /System/Library/PrivateFrameworks/AMPLibra |
| 391 | larisamor | 17 | 0 | 4266M | 4828 | ? | 0.2 | 0.1 | 0:27.98 | /System/Library/PrivateFrameworks/iTunesCl |
| 51040 | larisamor | 24 | 0 | 4197M | 3068 | R | 0.2 | 0.0 | 0:00.05 | htop |
| 369 | larisamor | 17 | 0 | 4537M | 6940 | ? | 0.2 | 0.1 | 5:28.80 | /System/Library/CoreServices/WiFiAgent.app |
| 532 | larisamor | 17 | 0 | 4582M | 8124 | ? | 0.1 | 0.1 | 1:49.54 | /System/Library/Frameworks/CoreServices.fr |
| 47093 | larisamor | 17 | 0 | 4370M | 25480 | ? | 0.1 | 0.3 | 4:26.21 | /System/Library/Frameworks/WebKit.framework |
| 442 | larisamor | 17 | 0 | 4258M | 5280 | ? | 0.0 | 0.1 | 2:19.99 | /System/Library/CoreServices/cloudpaired |
| 389 | larisamor | 8 | 0 | 4313M | 14468 | ? | 0.0 | 0.2 | 2:32.61 | /System/Library/PrivateFrameworks/Assistan |

The shell terminal

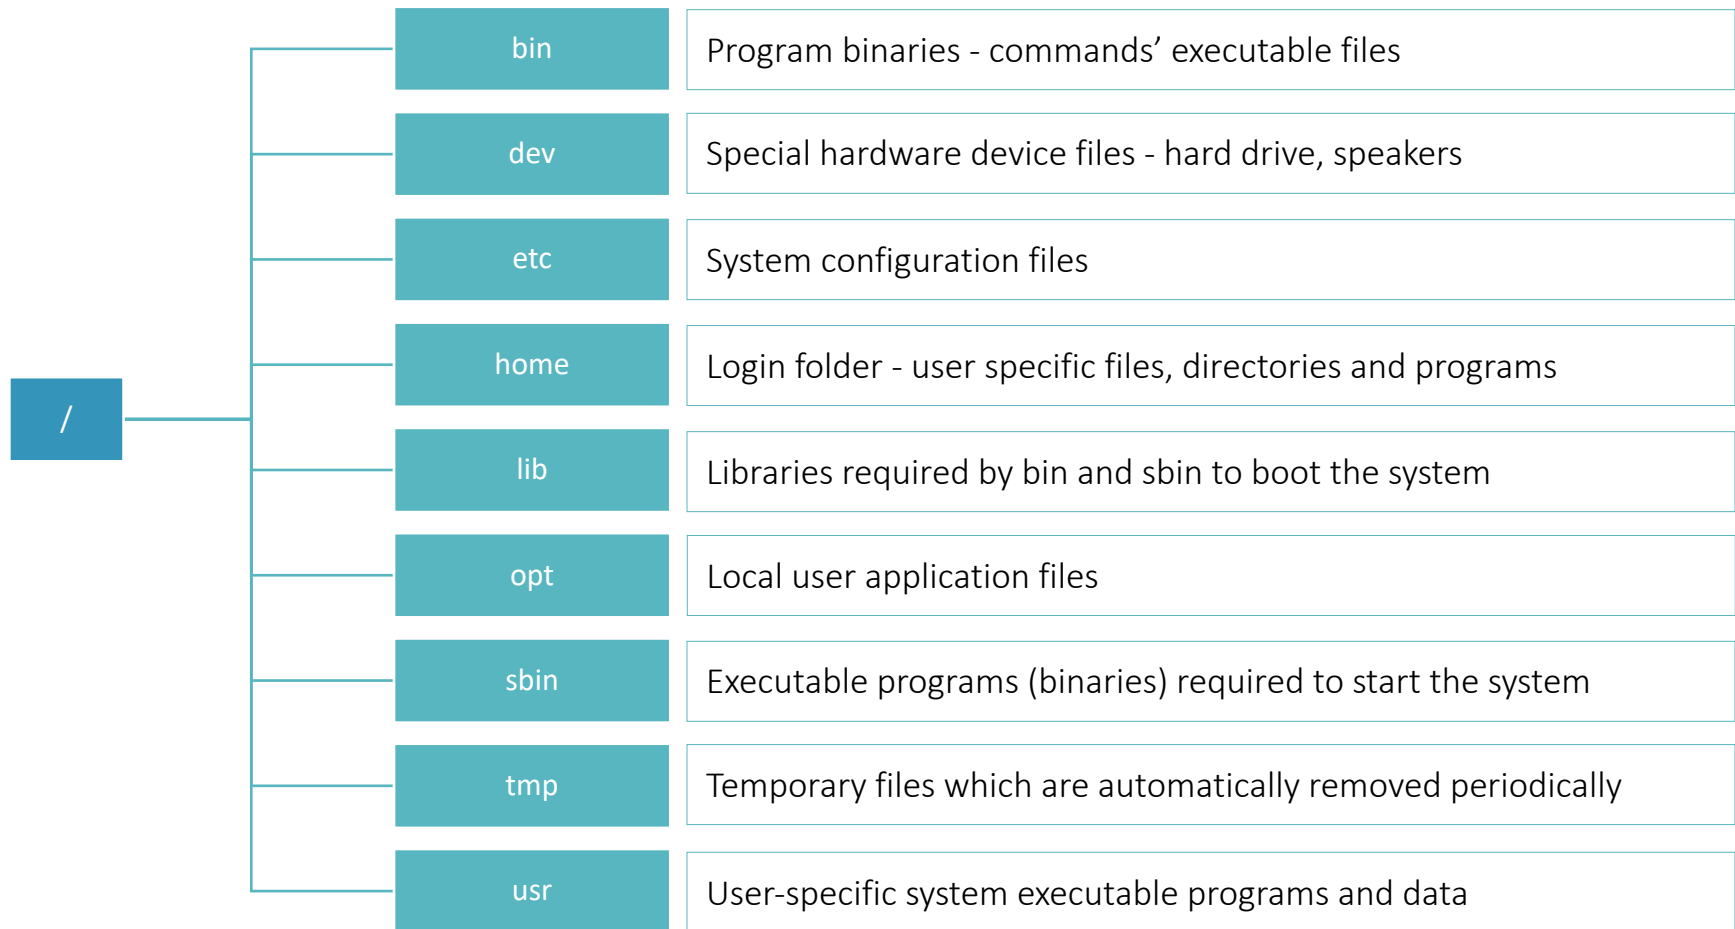


```
larisamoraless -zsh - 80x24
Last login: Wed Sep 15 18:11:33 on ttys003
(base) larisamoraless@MacBook-Pro-de-Larisa ~ % |
```


File system structure



File system structure



Hands on



1. Create a directory called folder1